

Tutorial

December 8, 2020

1 KwARG Tutorial

This tutorial gives an overview of the usage of KwARG and the types of output that it can produce.

To recreate the results on the command line, open Terminal and navigate to the directory where you have saved the binaries. Remove the ‘!’ from the commands below.

1.1 Available options

A list of the available options can be viewed by running KwARG with the ‘-h’, ‘-H’ or ‘-?’ option:

```
[1]: !./kwarg -h
```

Usage: ./kwarg [options] < [input]

The program reads data from the input file specified and greedily constructs a history with a low number of recombinations and recurrent mutations (homoplasies). The history is constructed by stepping backwards in time using coalescence, mutation and recombination events. At each point in this process, all possible next events (strictly speaking, only a useful subset of all possible next events) are considered, and the resulting ancestral states are scored. The scores are used to choose an event either at random or to proceed to an ancestral state with minimum score (see option -T). This process is NOT guaranteed to lead to a history with a minimum number of recombinations or the minimum number of homoplasies.

Legal options are:

- L[x] Provide an optional label x (should be an integer) to print at the start of each line.
- S[x] Specify cost of a sequencing error (default: x = 0.5).
- M[x] Specify cost of a recurrent mutation (default: x = 0.9).
- R[x] Specify cost of a single recombination (default: x = 1.0).
- C[x] Specify cost of two recombinations immediately following each other (default: x = 2.0).
- T[F] Set annealing temperature $0 < F < 700$ (default: $F = 30$). Scores are normalised to lie between 0 and 1, and the next step selected with probability proportional to $\exp(F * \text{normalised_score})$. Setting $F = 0$ corresponds to random selection among possible moves (not recommended). Setting F

= -1 corresponds to setting $F = \text{Infinity}$, and will force selection of the move with the minimum score at each step.

-V[x] If running a single iteration with given cost parameters, this controls the level of verbosity.

x = 0: no extra output

x = 1: during each neighbourhood search, output the number of neighbours explored, the move selected and its cost

x = 2: during each neighbourhood search, output the number of neighbours explored, the resulting configuration and cost of each neighbour, the move selected and its cost.

-b[name] Output a minimum recombination history to file name.

-d[name] Output ancestral recombination graph of minimum recombination history in dot format to file name.

-g[name] Output ancestral recombination graph of minimum recombination history in GDL format to file name.

-j[name] Output ancestral recombination graph of minimum recombination history in GML format to file name.

-t[name] Output list of marginal phylogenies for each site in Newick's 8:45 format to file name.

-D[name] Output list of marginal phylogenies for each site in dot format to file name.

-G[name] Output list of marginal phylogenies for each site in GDL format to file name.

-J[name] Output list of marginal phylogenies for each site in GML format to file name.

-I Marginal trees are only output one for each of the intervals between two recombination points, instead of one for each site.

-v[nodelabel] Use nodelabel convention for labelling nodes in ancestral recombination graphs. The possible conventions are: nodelabel =

- 'none': do not label nodes
- 'id': only label nodes representing sampled sequences, using their sequence ids from the data file, and nodes representing recombinations, indicating the recombination point
- 'sequence': label nodes with the inferred sequences; these sequences will be in binary format, with 0 representing wild type and 1 representing mutant type, even if the original data is not in binary format
- 'both': label nodes with both id and inferred sequence
- 'one': use only one label for a node, sequence id or inferred sequence

Default convention is id. The colour coding scheme used for the nodes is red for sequences in the input data, blue for recombination nodes, green for

- standard coalescent nodes, and yellow for the final coalescence into the most recent common ancestor.
- i Sequences not having a sequence id in the data file are assigned their index in the data file as id, e.g. the first sequence in the data file would be assigned '1' as id.
 - e Label edges in ancestral recombination graphs with the sites undergoing mutation along the edge.
 - o Assume input data is in own format. Default is to first try to parse data in own format, and if that fails to try to parse it in fasta format. Specifying this option, no attempt will be made to try to parse the data in fasta format.
 - f Assume input data is in fasta format. No attempt will be made to try to parse the data in own format. Note that the -o and the -f options override each other, so only the last one occurring in the command line will have an effect.
 - a Assume input consists of amino acid (protein) sequences using the one letter amino acid alphabet; anything not in the amino acid one letter alphabet is treated as an unresolved site (default is to assume sequences in binary, i.e. 0/1, format where anything but a 0 or a 1 is considered an unresolved site). If the most recent common ancestor is assumed known (see option -k), the first sequence in the input data is considered to specify the wild type of each site and is not included as part of the sample set.
 - n Assume input consists of nucleic sequences; anything but a/c/g/t/u is considered an unresolved site (default is to assume binary, i.e. 0/1, format where anything but a 0 or a 1 is considered an unresolved site). If the most recent common ancestor is assumed known (see option -k), the first sequence in the input data is considered to specify the wild type of each site and is not included as part of the sample set.
 - k Assume that the common ancestral sequence is known, i.e. that we know which is the wild type and which is the mutant in each site. If the data is in binary format, the all-0 sequence is assumed to be the common ancestral sequence (this does not need to be present in the data). If the data is in amino acid or nucleotide format, the common ancestral sequence has to be specified directly and is taken to be the first sequence in the data file (see options -a and -n)
 - Q[x] Sets the number of runs.
 - Z[x] Sets the random seed z (only one run is made in this case).
 - X[x] Provide an upper bound x on the number of recombinations needed for the input dataset (solutions with more than x recombinations will be abandoned).
 - s Turns off the header row of the results table.
 - h, -H -? Print this information and stop.

1.2 Running KwARG

We will first run KwARG on the Kreitman dataset, with 3 iterations for each of the default cost parameters:

```
[2]: !./kwarg -Q3 < kreitman_snp.txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
2601625704	30.0	-1.00	-1.00	1.00	2.00	0	0	7		153
0.01203100										
2833715205	30.0	-1.00	-1.00	1.00	2.00	NA	NA	NA		175
0.01934900										
2088249328	30.0	-1.00	-1.00	1.00	2.00	0	0	7		145
0.01337600										
3052951246	30.0	1.00	1.01	1.00	2.00	6	1	6		1444
0.57676100										
3353384182	30.0	1.00	1.01	1.00	2.00	4	1	4		1255
0.20606800										
3410089108	30.0	1.00	1.01	1.00	2.00	1	0	5		760
0.11520900										
1828302667	30.0	0.90	0.91	1.00	2.00	5	0	2		702
0.14626900										
3225485139	30.0	0.90	0.91	1.00	2.00	NA	NA	NA		1083
0.14772300										
3046884879	30.0	0.90	0.91	1.00	2.00	NA	NA	NA		849
0.09440700										
3198632652	30.0	0.80	0.81	1.00	2.00	NA	NA	NA		1157
0.21403300										
1986441690	30.0	0.80	0.81	1.00	2.00	3	0	4		914
0.10946100										
1773048665	30.0	0.80	0.81	1.00	2.00	2	1	4		780
0.11172300										
1838291634	30.0	0.70	0.71	1.00	2.00	NA	NA	NA		1167
0.15020400										
2143237569	30.0	0.70	0.71	1.00	2.00	NA	NA	NA		1086
0.15805000										
2303470289	30.0	0.70	0.71	1.00	2.00	NA	NA	NA		945
0.16666900										
2645221831	30.0	0.60	0.61	1.00	2.00	3	0	3		747
0.11292900										
3243300085	30.0	0.60	0.61	1.00	2.00	0	2	4		744
0.07948900										
1864865057	30.0	0.60	0.61	1.00	2.00	2	1	3		699
0.10499500										
2547814861	30.0	0.50	0.51	1.00	2.00	NA	NA	NA		942
0.09895000										
1921909335	30.0	0.50	0.51	1.00	2.00	3	0	3		822

0.11927000									
3185200706	30.0	0.50	0.51	1.00	2.00	5	0	2	820
0.10595400									
1627110433	30.0	0.40	0.41	1.00	2.00	NA	NA	NA	1010
0.22907700									
3056188657	30.0	0.40	0.41	1.00	2.00	NA	NA	NA	894
0.09132400									
1914316890	30.0	0.40	0.41	1.00	2.00	5	0	2	934
0.11731400									
3310285152	30.0	0.30	0.31	1.00	2.00	NA	NA	NA	931
0.18293100									
2830896468	30.0	0.30	0.31	1.00	2.00	5	0	2	715
0.09527100									
1765279040	30.0	0.30	0.31	1.00	2.00	5	0	2	723
0.10761400									
3393915342	30.0	0.20	0.21	1.00	2.00	5	0	2	723
0.10024600									
3187940384	30.0	0.20	0.21	1.00	2.00	NA	NA	NA	868
0.08212500									
2780186279	30.0	0.20	0.21	1.00	2.00	13	0	0	836
0.25446000									
2378498273	30.0	0.10	0.11	1.00	2.00	5	0	2	702
0.09103200									
1683804573	30.0	0.10	0.11	1.00	2.00	13	0	0	820
0.28792300									
2844859340	30.0	0.10	0.11	1.00	2.00	5	0	2	702
0.09583900									
3245102482	30.0	0.01	0.02	1.00	2.00	5	0	2	702
0.09950500									
1961747575	30.0	0.01	0.02	1.00	2.00	9	0	1	889
0.10874100									
2780771945	30.0	0.01	0.02	1.00	2.00	7	5	0	736
0.09567800									
2276621186	30.0	1.00	1.10	-1.00	-1.00	NA	NA	NA	1169
0.21339400									
2917570376	30.0	1.00	1.10	-1.00	-1.00	NA	NA	NA	1432
0.20109900									
2549774585	30.0	1.00	1.10	-1.00	-1.00	NA	NA	NA	1589
0.25289100									

The output table gives: - Seed: the random seed needed to rerun this particular iteration (demonstrated below) - Temp: the annealing temperature used (default: 30) - SE_cost, RM_cost, R_cost, RR_cost: the cost parameter for a sequencing error, recurrent mutation, single recombination, and two consecutive recombination events, respectively - SE, RM, R: the number of each type of event in the solution - N_states: the total number of states considered when constructing one-step neighbourhoods - Time: CPU time for each iteration.

Lines where number of events is shown as 'NA' correspond to runs which were identified as sub-optimal and terminated before completion.

KwARG can also be run with specified costs, as follows:

```
[3]: !./kwarg -S0.2,0.4 -M0.3,0.5 -Q5 < kreitman_snp.txt
```

Time	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
2611612860	30.0	0.20	0.30	1.00	2.00	8	2	0		715
0.09174800										
2186755964	30.0	0.20	0.30	1.00	2.00	6	0	2		813
0.12789900										
3641081466	30.0	0.20	0.30	1.00	2.00	NA	NA	NA		1297
0.13407100										
2458405652	30.0	0.20	0.30	1.00	2.00	5	0	2		702
0.10066300										
2963344908	30.0	0.20	0.30	1.00	2.00	NA	NA	NA		937
0.11827100										
2681191801	30.0	0.40	0.50	1.00	2.00	4	1	2		770
0.08341100										
2159394232	30.0	0.40	0.50	1.00	2.00	NA	NA	NA		1088
0.12439800										
2526916363	30.0	0.40	0.50	1.00	2.00	NA	NA	NA		814
0.29432900										
2930465543	30.0	0.40	0.50	1.00	2.00	NA	NA	NA		792
0.11778200										
2775244447	30.0	0.40	0.50	1.00	2.00	NA	NA	NA		1038
0.09858000										

The same number of SE_cost and RM_cost parameters should be specified, separated by commas.

The options ‘-R’ and ‘-C’ can be used to specify the cost of single and double recombination events, this defaults to 1.0 and 2.0, respectively. If these options are used, then SE_cost and RM_cost must also be provided.

1.2.1 Input formats

Binary data The input data can be in 0/1 binary format, and any other symbols will be interpreted as missing data.

```
[4]: !cat example_data_1.txt
```

```
00010010
1001101x
11001x01
01100101
```

Sequence and site labels can also be provided, as follows:

```
[5]: !cat example_data_2.txt
```

```
#> Seq1
00010010
#> Seq2
1001101-
#> Seq3
11001-01
#> Seq4
01100101
#positions: 2 3 7 9 10 11 13 14
```

FASTA format Input data in fasta format can be provided, for instance:

```
[6]: !cat example_data_3.fasta
```

```
>Seq1
ACTTAAGG
>Seq2
TCTTTAGN
>Seq3
TGTATNAA
>Seq4
AGCAATAA
```

In this case, to denote that the data is provided in fasta format in nucleotide representation, the ‘-f’ and ‘-n’ flags should be used:

```
[7]: !./kwarg -f -n -S0.5 -M0.9 < example_data_3.fasta
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
	2616606438	30.0	0.50	0.90	1.00	2.00	2	0	0	36
	0.00114700									

1.2.2 Annealing parameter

We can change the annealing temperature using the ‘-T’ option: - ‘-T30’ is the default temperature of 30. - ‘-T0’ means the next step is chosen uniformly at random among all available moves (this is not recommended). - ‘-T-1’ signifies $T = \infty$, so the next step is chosen among all available moves with the minimum score.

```
[8]: !./kwarg -T-1 -Q3 < kreitman_snp.txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
	2616606438	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	146
	0.01183900									
	1863276342	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	146
	0.01072300									

1634361799	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	167
0.01147600									
2120612875	-1.0	1.00	1.01	1.00	2.00	1	0	6	797
0.10581000									
2768325270	-1.0	1.00	1.01	1.00	2.00	1	0	6	797
0.09598300									
3390484970	-1.0	1.00	1.01	1.00	2.00	1	0	6	797
0.09088400									
3398681837	-1.0	0.90	0.91	1.00	2.00	3	0	3	773
0.09291300									
2177631974	-1.0	0.90	0.91	1.00	2.00	3	0	3	773
0.10779600									
2872255873	-1.0	0.90	0.91	1.00	2.00	3	0	3	773
0.10341900									
3637292166	-1.0	0.80	0.81	1.00	2.00	3	0	3	789
0.09916600									
2229965343	-1.0	0.80	0.81	1.00	2.00	3	0	3	789
0.09505000									
2546633691	-1.0	0.80	0.81	1.00	2.00	3	0	3	789
0.09396800									
2918284069	-1.0	0.70	0.71	1.00	2.00	3	0	3	789
0.09800800									
2944500523	-1.0	0.70	0.71	1.00	2.00	3	0	3	789
0.09861100									
1872404186	-1.0	0.70	0.71	1.00	2.00	3	0	3	789
0.10707600									
1710428850	-1.0	0.60	0.61	1.00	2.00	3	0	3	789
0.09456800									
1991187006	-1.0	0.60	0.61	1.00	2.00	3	0	3	789
0.09400000									
2011993707	-1.0	0.60	0.61	1.00	2.00	3	0	3	789
0.09674300									
2817200563	-1.0	0.50	0.51	1.00	2.00	5	0	2	791
0.10336500									
2265418323	-1.0	0.50	0.51	1.00	2.00	5	0	2	791
0.11140600									
2065884106	-1.0	0.50	0.51	1.00	2.00	5	0	2	791
0.10161900									
2194177202	-1.0	0.40	0.41	1.00	2.00	5	0	2	723
0.13337700									
3484274397	-1.0	0.40	0.41	1.00	2.00	5	0	2	723
0.11459100									
2622644005	-1.0	0.40	0.41	1.00	2.00	5	0	2	723
0.12060700									
2837983743	-1.0	0.30	0.31	1.00	2.00	5	0	2	723
0.10573600									
2572373776	-1.0	0.30	0.31	1.00	2.00	5	0	2	723
0.10198800									

3485362873	-1.0	0.30	0.31	1.00	2.00	5	0	2	723
0.10146700									
1941959837	-1.0	0.20	0.21	1.00	2.00	5	0	2	723
0.10453900									
3690368037	-1.0	0.20	0.21	1.00	2.00	5	0	2	723
0.10957000									
3011058638	-1.0	0.20	0.21	1.00	2.00	5	0	2	723
0.10870000									
2981580430	-1.0	0.10	0.11	1.00	2.00	5	0	2	702
0.10073700									
3621259533	-1.0	0.10	0.11	1.00	2.00	5	0	2	702
0.09977600									
2285808715	-1.0	0.10	0.11	1.00	2.00	5	0	2	702
0.09485400									
3548241670	-1.0	0.01	0.02	1.00	2.00	5	0	2	702
0.09851200									
3359367899	-1.0	0.01	0.02	1.00	2.00	5	0	2	702
0.10339400									
3409163983	-1.0	0.01	0.02	1.00	2.00	5	0	2	702
0.10292400									
1662077177	-1.0	1.00	1.10	-1.00	-1.00	16	0	0	724
0.11880500									
2811479899	-1.0	1.00	1.10	-1.00	-1.00	15	0	0	716
0.11003600									
2016335473	-1.0	1.00	1.10	-1.00	-1.00	15	0	0	716
0.10752000									

Several values of the annealing parameter can be provided, separated by commas: for instance ‘-T10,20,50’.

1.2.3 Turning off recurrent mutations

Specifying ‘-S-1 -M-1’ turns off recurrent mutations, so in this case an upper bound on Rmin is computed under the infinite sites assumption:

```
[9]: !./kwarg -S-1 -M-1 -T-1 -Q5 < kreitman_snp.txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
2623264542	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	146	
0.01144100										
2147798063	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	167	
0.01107900										
1986743730	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	146	
0.01033200										
1996696503	-1.0	-1.00	-1.00	1.00	2.00	0	0	7	146	
0.01076200										

```

2947321370  -1.0    -1.00   -1.00    1.00    2.00   0   0   7      167
0.01370200

```

1.2.4 Turning off recombination

Specifying ‘-R-1 -C-1’ turns off the possibility of recombinations, so an upper bound on the minimum parsimony score is computed:

```
[10]: !./kwarg -S1.0 -M1.1 -R-1 -C-1 -Q5 < kreitman_snp.txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
2623264542	30.0	1.00	1.10	-1.00	-1.00	18	8	0		2018
0.43616500										
1996696503	30.0	1.00	1.10	-1.00	-1.00	NA	NA	NA		2672
0.68387300										
2947321371	30.0	1.00	1.10	-1.00	-1.00	24	2	0		1875
1.12214200										
2838510949	30.0	1.00	1.10	-1.00	-1.00	20	1	0		1683
0.34656700										
1792471208	30.0	1.00	1.10	-1.00	-1.00	NA	NA	NA		1691
0.43594300										

1.2.5 Specifying the ancestral (root) sequence

It is possible to specify a particular sequence as ancestral to the sample (corresponding to the root of the ARG), using the ‘-k’ option.

If the input data is in binary format, then the all-zero sequence will be assumed to be ancestral (whether or not this is included in the data).

If the input data is in nucleotide or amino acid representation, then the first sequence in the data will be taken as ancestral.

The effect on the resulting history is illustrated further below.

1.3 Outputs

We can re-run a particular instance via specifying the random seed, and output the history as text. Here we are choosing seed Z3391369848, SE_cost=0.01 and RM_cost=0.02. We expect to see 5 sequencing errors, 0 recurrent mutations and 2 recombinations in the output.

The “-b” option will save the history in file “example_history.txt”.

The “-d” option will output a picture of the corresponding ARG in DOT format and save it in the file “example_arg.dot”. The “-e” options tells KwARG to label the edges with mutations.

If an output option is specified, an output file name must also be given.

```
[11]: !./kwarg -S0.01 -M0.02 -Z3391369848 -bexample_history.txt -dexample_arg.dot -e
      ↪ kreitman_snp.txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time										
	3391369848	30.0	0.01	0.02	1.00	2.00	5	0	2	702
	0.12403500									

```
[12]: !cat example_history.txt
```

```
Mutation of site 6 in sequence 11
Mutation of site 7 in sequence 11
Mutation of site 8 in sequence 11
Mutation of site 10 in sequence 1
Mutation of site 14 in sequence 11
Mutation of site 15 in sequence 5
Mutation of site 21 in sequence 11
Mutation of site 25 in sequence 11
Mutation of site 39 in sequence 4
Mutation of site 40 in sequence 4
Mutation of site 41 in sequence 3
Mutation of site 42 in sequence 6
Mutation of site 43 in sequence 3
Coalescing sequences 8 and 9
Coalescing sequences 8 and 10
Mutation of site 13 in sequence 8
Mutation of site 38 in sequence 8
---->Sequencing error at site 36 in sequence 6
---->Stretch of sequencing errors spanning 2 sites:
---->Sequencing error at site 17 in sequence 5
---->Sequencing error at site 18 in sequence 5
Mutation of site 17 in sequence 4
Mutation of site 18 in sequence 4
Coalescing sequences 3 and 4
Mutation of site 36 in sequence 3
---->Sequencing error at site 9 in sequence 1
---->Sequencing error at site 3 in sequence 2
Coalescing sequences 1 and 2
Mutation of site 19 in sequence 1
Mutation of site 20 in sequence 1
Mutation of site 22 in sequence 1
Mutation of site 23 in sequence 1
Mutation of site 24 in sequence 1
Mutation of site 26 in sequence 1
Mutation of site 27 in sequence 1
Mutation of site 28 in sequence 1
Mutation of site 29 in sequence 1
Coalescing sequences 1 and 3
```


```

Mutation of site 16 in sequence 1
Mutation of site 37 in sequence 1
---->Recombination in sequence 5 after site 29; suffix is new sequence 12
Coalescing sequences 1 and 12
Mutation of site 31 in sequence 1
Mutation of site 33 in sequence 1
Mutation of site 35 in sequence 1
---->Recombination in sequence 5 after site 3; prefix is new sequence 13
Coalescing sequences 1 and 13
Mutation of site 3 in sequence 1
Coalescing sequences 8 and 5
Coalescing sequences 1 and 6
Mutation of site 11 in sequence 8
Mutation of site 12 in sequence 8
Mutation of site 30 in sequence 1
Mutation of site 32 in sequence 1
Mutation of site 34 in sequence 1
Coalescing sequences 1 and 7
Coalescing sequences 8 and 11
Mutation of site 1 in sequence 8
Mutation of site 2 in sequence 8
Mutation of site 4 in sequence 8
Mutation of site 5 in sequence 8
Mutation of site 9 in sequence 1
Coalescing sequences 1 and 8
Total: 5 sequencing errors, 0 recurrent mutations, 2 recombinations.

```

If GraphViz is installed, it can be used to convert the DOT file to a png.

```
[13]: !dot -Tpng -o example_arg.png example_arg.dot
```



Recurrent mutations are labelled with '*'. Recombination nodes are coloured blue, with the number inside the node denoting the recombination breakpoint. The parts of the genome to the left and right of the breakpoint are inherited from two different parent nodes. The edges leading to these nodes are labelled 'P' and 'S', which stands for 'prefix' and 'suffix', respectively.

See the help for a full list of possible output formats. For instance, we can also print out the local trees for each interval in Newick format as follows:

```
[14]: !./kward -S0.01 -M0.02 -Z3257491408 -texample_local_trees.txt -I < kreitman_snp.
      ↪txt
```

	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R	N_states
Time	3257491408	30.0	0.01	0.02	1.00	2.00	5	0	2	702
										0.16322600

The '-I' option means a tree is given for each interval between recombination breakpoints. Not specifying this option will mean that one tree is produced for each site.

```
[15]: !cat example_local_trees.txt
```

```
(((((1,2),(3,4)),5),6),7),((8,9),10),11))1-3;
(((1,2),(3,4)),6),7),((5,((8,9),10)),11))4-29;
((((1,2),(3,4)),5),6),7),((8,9),10),11))30-43;
```

1.3.1 Specifying the ancestral sequence

Compare the output ARGs when the ancestral sequence is and is not specified:


```
[16]: !cat example_data_2.txt
```


```
#> Seq1
00010010
#> Seq2
1001101-
#> Seq3
11001-01
#> Seq4
01100101
#positions: 2 3 7 9 10 11 13 14
```

```
[17]: !./kwarg -L1 -Z3066074262 -dexample_ancestral.dot -k -e -vboth < example_data_2.
      ↪txt
      !./kwarg -L2 -Z3066074262 -dexample_nonancestral.dot -e -vboth -s <
      ↪example_data_2.txt
      !dot -Tpng -o example_ancestral_arg.png example_ancestral.dot
      !dot -Tpng -o example_nonancestral_arg.png example_nonancestral.dot
```

	Ref	Seed	Temp	SE_cost	RM_cost	R_cost	RR_cost	SE	RM	R
N_states		Time								
	1	3066074262	30.0	0.50	0.90	1.00	2.00	2	0	0
55	0.01355600									
	2	3066074262	30.0	0.50	0.90	1.00	2.00	2	0	0
36	0.00166900									

Here '-vboth' means that the leaves are marked with their corresponding reference (if specified), and all nodes are labelled with the corresponding sequence.





1.4 Simplify

This is a small program which reduces the input dataset using the ‘Clean’ algorithm, removing all mutations and coalescing all possible sequences until no further reduction is possible. The input data types are the same as for KwARG.

```
[18]: !./simplify < kreitman_snp.txt
```

```

Input dataset: 11 sequences, 43 sites
Reduced dataset: 9 sequences, 16 sites
0100010101010101
0001010101010101
0101010001010111
0101011001010111
0110101001010100
0001000000000010
0001000010101000
1010100010101000
1010000010101000
Sequences:
X 1 2 3 4 5 6 X 10
Sites:
X 2 X 8 X 15 X X 29 30 31 32 33 34 35 36

```

The output sequence labels show either the number of the sequence in the input dataset, or ‘X’ if the sequence in the output has been coalesced with another. The site labels show either the number of the site as in the input dataset, or ‘X’ if the column corresponds to multiple sites collapsed into one.

1.5 Flip

This is a small program which flips the nucleotides at the given sequence and site (from 0 to 1) and outputs the resulting amended dataset.

```
[19]: !cat kreitman_snp.txt
```

```

0010000001000001001101110111101010101000000
0000000010000001001101110111101010101000000
00100000100000010000000000000001010111000101
00100000100000011100000000000001010111011000
00111000001100101100000000000001010100000000
000000001000000000000000000000000000000000010000010
00000000100000000000000000000000010101000000000
110110000011100000000000000000010101000100000
110110000011100000000000000000010101000100000
110110000011100000000000000000010101000100000
1101111100000100000010001000010101000000000

```

[20]: `!./flip -q6,8 -s37,1 < kreitman_snp.txt`

```

0010000001000001001101110111101010101000000
0000000010000001001101110111101010101000000
00100000100000010000000000000001010111000101
00100000100000011100000000000001010111011000
00111000001100101100000000000001010100000000
00000000100000000000000000000000000000000011000010
0000000010000000000000000000000001010100000000
010110000011100000000000000000010101000100000
110110000011100000000000000000010101000100000
110110000011100000000000000000010101000100000
1101111100000100000010001000010101000000000

```

Using the option ‘-bfilename.txt’ will save the resulting dataset to filename.txt.